PHREAKNIC 27

Strategies for the Post-LLM Job Market

It's Always Risk Management

David Cheeseman

Weighing what could go wrong against what you stand to gain

Disclaimer

Never trust the advice of a single person, including me.

Opinions are my own and do not represent my current or past employers.

Outline

1. **Data science primer** — how models are trained, and how they fail

2. **How LLMs have changed work** — everyone is a risk manager and a mini project lead

3. **How LLMs have changed the job search** — the hiring pipeline and your own toolkit

4. **Tactics** — for everyone, and by what you want from the job

5. **Conclusion** — what changed, and what has not

How a Model Gets Built and Graded

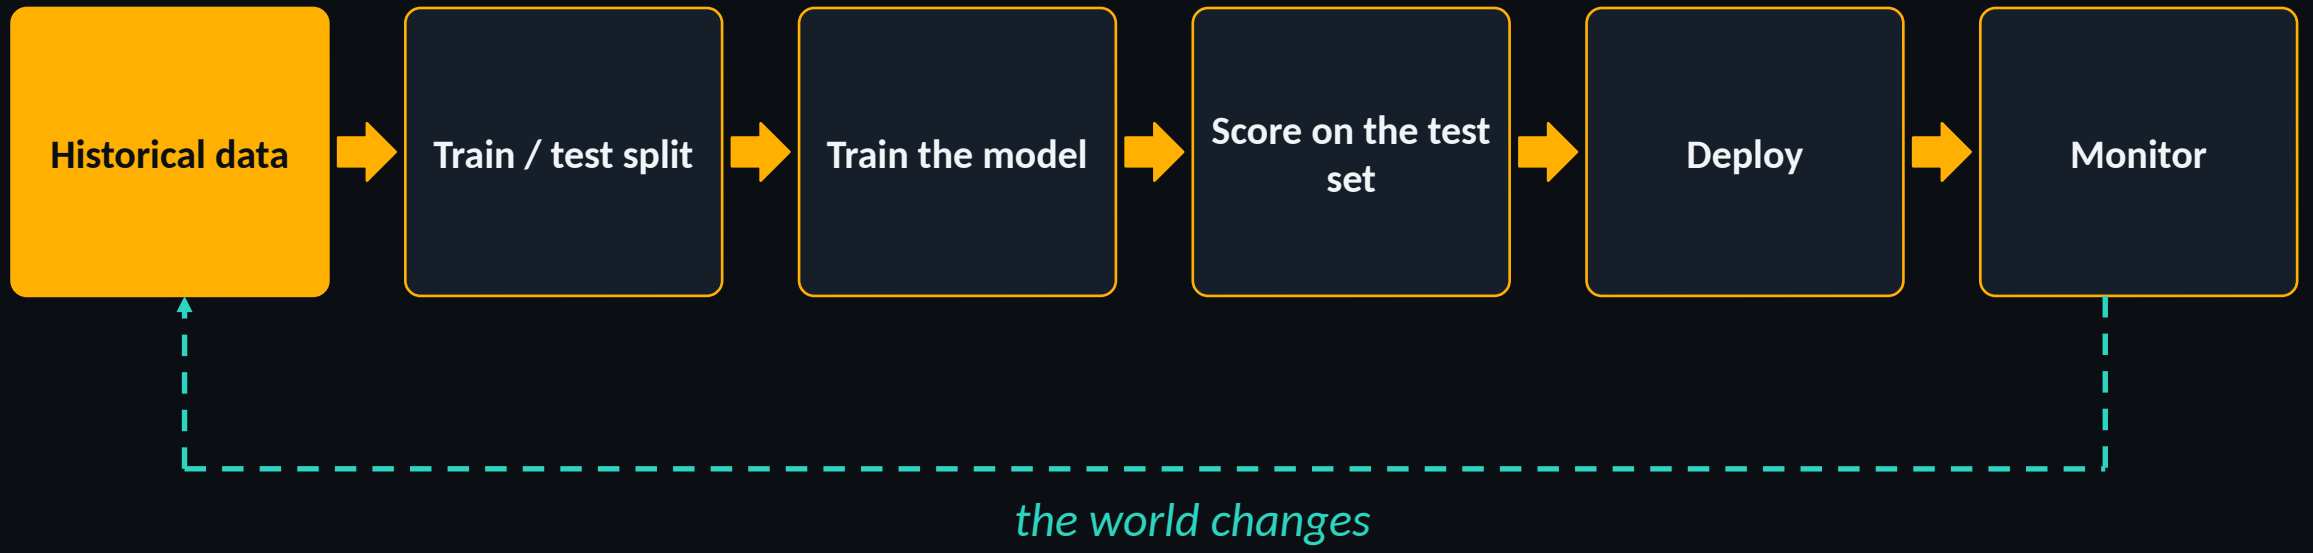

A model learns the past. If you hired badly before, it learns that too.

How LLMs Are Trained

1

Pretraining

Predict the next token across a huge corpus of text and code

2

Fine-tuning

Curated examples teach the shape of a helpful answer

3

Preference tuning

Raters (human or model) rank outputs; the model drifts toward the winners

4

Deployment

System prompt, tools, and context steer it; a sampler picks every token

Traditional AI: one task, labeled data, a test set with a score.

LLM: general text prediction, then shaped to sound helpful.

It’s impossible to check every feature input/prompt context so the probability predictions/prompts are wrong can be reduced but not eliminated.

Why the Same Prompt Gives Different Answers

Prompt: "The best database for this app is ..."

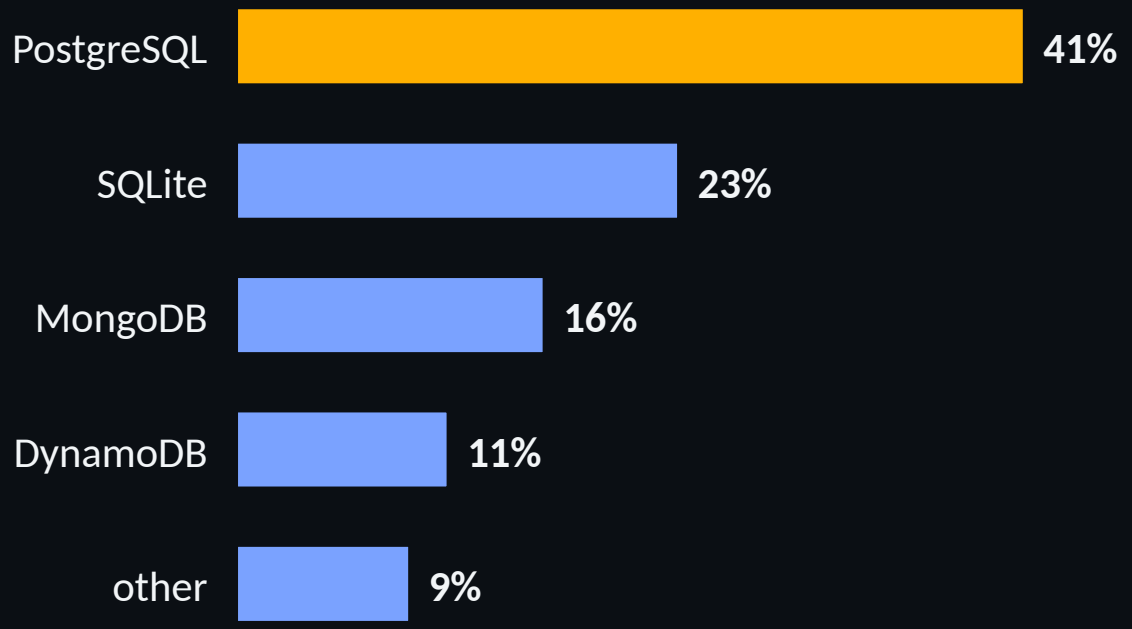

Illustrative numbers, not measured.

WHAT HAPPENS

- The model scores every possible next token
- A sampler draws one; temperature widens or narrows the draw
  - You can set the temperature to 0 to get the same answer for every prompt
- Every draw changes what follows
- Run it again and the draw differs

Same prompt, same model, different project.

Underfit and Overfit

Demo footage: a driving AI learning from a human driver (GTA V). Plays automatically, muted, loops.

UNDERFIT

Too simple: misses the pattern

In hiring: a keyword filter rejects "Kubernetes" because the posting said "K8s".

GOOD FIT

Captures the pattern

Holds up on candidates it has never seen.

OVERFIT

Memorized the examples

In hiring: "every good hire we ever had went to these three schools."

Hallucination

Demo footage: pedestrian avoidance (GTA V). Plays automatically, muted, loops.

Confident, fluent, wrong.

WHERE THE WORD CAME FROM

2000	**Face hallucination**	A feature: inventing plausible detail to sharpen a face
2018	**Translation and captioning**	An error: output the input doesn't support
2020	**Intrinsic vs. extrinsic**	Contradicts the source, or adds what can't be verified
2022+	**The ChatGPT era**	Any confident false statement

HOT TAKE: Hallucination should be called Erros

HOT TAKE: You NEED to look at the code!

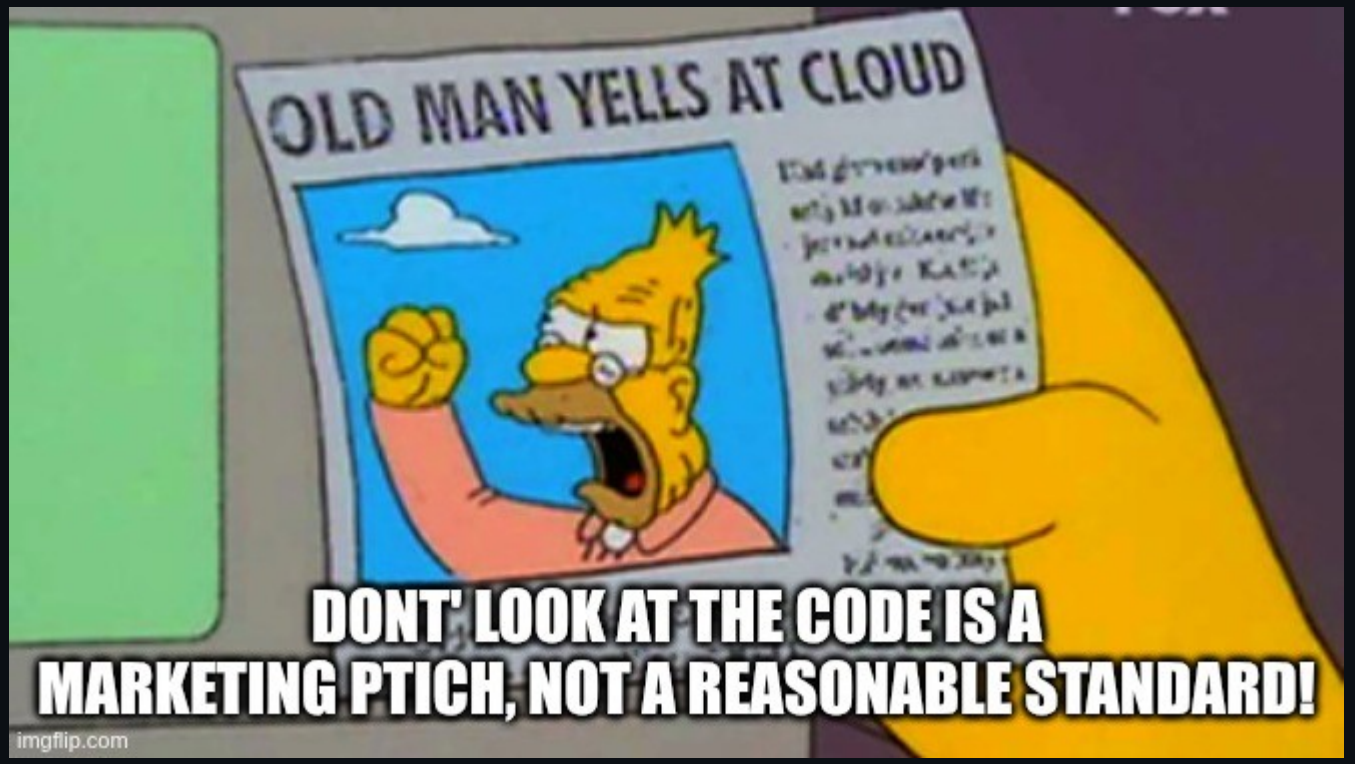

Everyone Is Now a Risk Manager

Decision	Before	Now
Design	An architect reviews the approach; weeks to decide	Whether to accept the vibe coded solution as a PoC or ready for prod manage SOLUTION EXPLOSION
Data management	A data steward decides retention, access, and where data may go	Make decision about whether AI vendors can be trusted with real data or if you have to synthesize
Security	A security team reviews auth, crypto, and dependencies	Decide on security controls for dependencies and CWE's (disabling TLS, old/vulnerable dependencies)

Everyone can now work much faster with less technical knowledge, but the risks of developing software has not changed.

Risk = Probability × Impact

	Event	Probability	Impact
pic	Foot Lettuce	Medium	Medium
pic	It's always DNS	High	High
pic	Stuck on a camera filter in court	Low	Medium
pic	The USB plug is upside down	High	Low

ISO 31000: risk is the "effect of uncertainty on objectives", positive or negative.

Qualitative vs. Quantitative

QUALITATIVE

Low | Medium | High

What you use when you don't have data. Most job seekers live here, and that's fine.

QUANTITATIVE: THE LEDGER

0.02 × $40,000 = $800

expected loss

0.6 × $10,000 = $6,000

expected gain

Both are valid; know which one you're doing, and always fill in both lines.

Levels of LLM Use in Organizations

Level	Example	Risk	Benefit
1. Periodic use	An engineer pastes a stack trace into a chat window	Data leaves the building; no record, no review	Cheap and fast; nothing in production depends on it
2. Coded, deterministic solutions	The LLM writes a log parser; it is reviewed, tested, and runs with no LLM	Subtle bugs and bad dependencies pass an unread review	Repeatable, testable, and no cost per run
3. LLMs in the product (non-deterministic)	A support bot or resume screener calls a model on every request	Wrong, biased, or leaky answers at scale; per-call cost; vendor drift	Handles messy input no rule set could
4. Fully automated development and deployment	An agent takes a ticket, writes the code, merges it, and ships it	Nobody understands the system; runaway spend; unreviewed changes in production	Throughput; one person ships a solution that operates at team levels of productivity

Each step up trades predictability for productivity. Each organization will have a different risk appetite in general and it may vary solution to solution.

Everyone Is a Mini Project Lead

Illustrative runs.

Vague requirements get a different project every time.

Plan the Project, Not Just the Prompt

Project structure

Where code, tests, and docs live

Language and versions

Pinned, so each run starts the same

Dependency policy

What may be added, and who approves it

Coding standard

One style, enforced by tools

Tests

What counts as proof it works

The gate

One command that must pass before done

Security and data rules

What may never leave the repo or the network

Decisions and docs

Written down so the agent can read them

HAVE THE MODEL INTERVIEW YOU TO MAKE THESE DECISIONS AND DOCUMENT IN A DESIGN DOCUMENT!

Code Quality Tools: Checking the Code

Tool class (example)	Benefit to the project	Why LLM-written code gets more predictable
Linters (Ruff, ESLint)	Catch bugs and bad patterns early	Turn style rules into pass/fail feedback the agent can read and fix
Formatters (Prettier, Black)	One style; clean diffs	Remove formatting variance between runs; reviewers see only logic
Security / CWE scanners (Sonarqube Semgrep, Bandit, CodeQL)	Find injection, unsafe APIs, hardcoded secrets	Flag known weakness classes the model reproduces from its training data
Dependency checkers (OSV-Scanner, pip-audit)	Catch vulnerable, abandoned, or nonexistent packages	Reject package names the model invented
Test suites (pytest, Vitest)	Prove behavior; protect refactors	Act as the spec; the agent iterates until they pass instead of declaring victory

Code Quality Tools: Mapping and Pruning the Code

Tool class (example)	Benefit to the project	Why LLM-written code gets more predictable
Auto-documentation (TypeDoc, Sphinx)	Docs stay in sync with code	Forces typed, documented interfaces; stale docs fail the build
Code mappers (Graphlify)	A map of modules and call paths	Gives the agent structure to follow instead of guessing which file owns what, reduces token costs
Dead code finders (Vulture, knip)	Remove unused code and dependencies	Catch what agents leave behind when they try several approaches, reduces test bloat and execution time
Profilers (py-spy, cProfile)	Show where time and memory really go	Replace 'looks fast' with a number the agent must beat

Consolidate into one or more gate commands (Makefile FTW) and/or give visibility into builds for your agents.

Instruction Files, Agents, and Skills

Instruction files

AGENTS.md, CLAUDE.md

- DESIGN.md
  - High level design
- README.md
  - End User Only
- CONTRIBUTING.md
  - Human AND agent guidance
- AGENTS.md
  - Agent SPECIFIC guidance
- CLAUDE.md/copitlot-instructions.md
  - Pointer to AGENTS.md ONLY

Agents and subagents

narrow, scoped helpers

- One focused job each
- A description that says when to delegate
- Allowlist the tools
- Tell a reviewer what to skip

Skills

procedures loaded on demand

- Say what it does and when to use it
- Body under 500 lines
- Detail one level down
- Test it against real tasks

HOT TAKE: The Sr and Jr Developer Agent Trap

THE TRAP

An agent whose job is to be worse than the model can be.

SEPARATE BY

- Task: writer and reviewer, not senior and junior
- Tools: each gets only what it needs
- Context: independent jobs, not shared guesses
- Judges: tests and scanners, not a persona

Engineer your model selection NOT your persona skill levels!

Best Practices That Contradict Each Other

Question	One position	The opposing position
Does AI make developers faster?	Developers believed they were about 20% faster (METR 2025)	They measured 19% slower; 2026 update inconclusive
Does it improve delivery?	DORA 2025: AI amplifies strong teams	DORA 2024: -1.5% throughput, -7.2% stability
Generate or hand-write AGENTS.md?	Docs: give the agent what it cannot see	2026 study: generated context files lowered success, raised cost 20%+
Many agents or one?	Anthropic: +90.2% on a research eval at ~15x tokens	Cognition: partial context makes conflicting decisions
Is the code better or just more?	Vendors: faster delivery of features	GitClear: clones up 8.3% to 12.3%, refactoring down

Sources in notes and on the References slide. Several figures are vendor or survey estimates.

LLM Solution Cost Control

Tokens

cost per call × calls

Errors

false positive vs. false negative price tags

Agents

triggers, loops, context size

What triggers the agent	Who controls the rate	Worst case
You (Skill to design a database schema)	You	You forget it's running
A schedule (Summarize a daily log)	Whoever set the cron	Runs 24/7 on nothing
An external user (Phishing filter)	Not you	Denial of Service/Wallet Vector
A loop inside the flow (Incident triage)	The prompt and the stop condition	**Unbounded**

Always consider if what you're getting the agent flow to do could be accomplished by asking the agent to write a script instead.

Cost: Redesign because of wrong platform/compliance

Month 1 Prototype summarizer built on a commercial chat API

Month 3 Real contract cost data goes in; demo wows leadership

Month 5 Compliance asks: is this service FedRAMP Moderate or equivalent?

Month 6 Rebuild on an authorized platform; sunk cost written off

THE CHECK THAT WAS SKIPPED

- Does the contract mark this data as CUI?
- Is the service FedRAMP Moderate or equivalent (DFARS 7012)?
- Does it meet NIST SP 800-171 controls?
- Who signs off before real data goes in?

Constructed example. Public financial reports are not CUI; contract-identified non-public data can be.

Consider ALL stakeholders! Users, Data Owners, Security, Governance, etc.

Skill Atrophy

If the AI stops working or you run out of tokens tomorrow, can you still do the job?

WHAT THE STUDIES SAY

- Anthropic trial (52 devs, 2026): quiz 50% with AI vs 67% without; biggest gap on debugging
- Delegating scored ~39%; generate-then-understand ~86%
- CHI 2025 (319 workers): more trust in AI, less critical thinking

KEEP THE MUSCLE

- Do some work unassisted, on purpose
- Ask the model to explain before you accept
- Keep a runbook and a fallback for outages
- Be able to debug what it wrote

Small samples; one is self-reported. Treat as evidence of risk, not proof of loss.

Your Company Will Not Fire Claude

Who	Held accountable?	What happens
The model	No	Nothing. It has no job, no license, and no testimony.
You	Yes	Your company fires you; a licensed professional can be sanctioned
Your C-suite	Yes	They are called before Congress, and they will still find someone to blame

The model does not take the blame. You do.

Poll: Why Are You Looking?

LIVE POLL

I need a job now

More pay

Better job satisfaction

Chasing the dream job

Not looking (yet)

Know Your Goal

Need it now

PAYOFF

stability

RISK

financial risk

More pay

PAYOFF

money

RISK

opportunity cost

Satisfaction

PAYOFF

fit

RISK

fit risk

Dream job

PAYOFF

the dream

RISK

long odds

Risk/Benefits of using LLMs in the Job Search

Hiring pipeline

BENEFIT

Helps them sift hundreds of applicants

RISK

A bad filter can drop you

Resume writing

BENEFIT

Catches errors, fits the filters

RISK

Everyone sounds the same; embellishment

Open source intel

BENEFIT

Learn what the company really does

RISK

They can do the same to you

Mock interviews

BENEFIT

Find your weak spots out loud

RISK

None. Just do this.

What the HR Screener Sees

One remote, entry-level Python opening

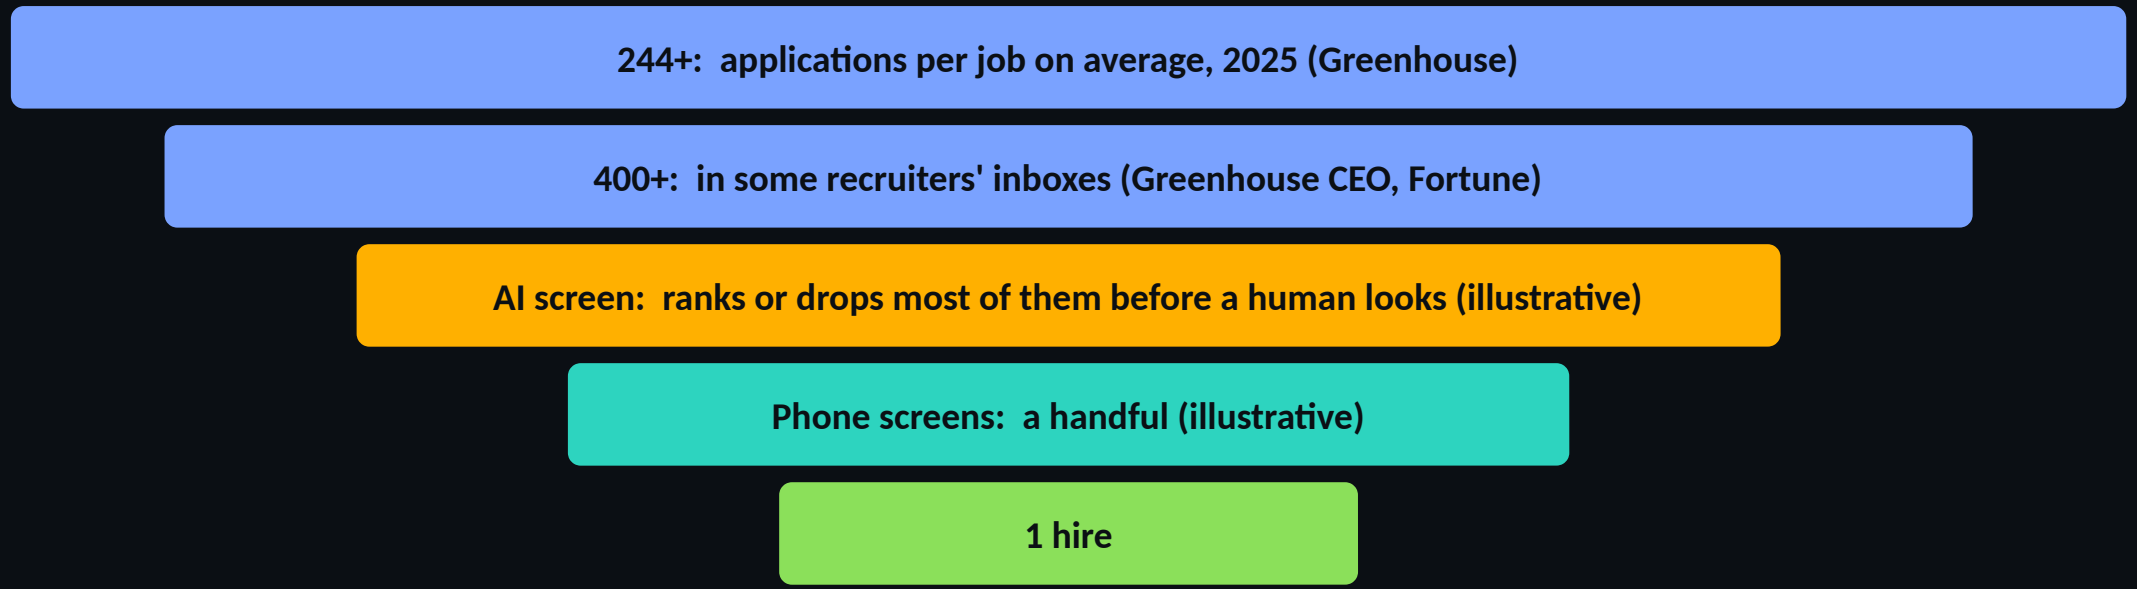

Applicant counts are sourced; the funnel's lower stages are illustrative.

BENEFIT

Helps a person sift hundreds of applicants.

RISK

A model, not a person, makes the first cut on you.

Five Products As Examples

Product	What employers get	What happened	Lesson
Workday + HiredScore	A ranked shortlist inside their ATS	Mobley v. Workday: bias claims, vendor-as-agent theory proceeds	You can't fully transfer risk to a vendor
Paradox "Olivia" (Workday since Oct 2025)	24/7 screening and scheduling	McHire admin login: 123456 / 123456, up to 64M records exposed	Your resume is only as safe as their worst login
LinkedIn Hiring Assistant	Agent that sources and screens	62% fewer profiles reviewed, per LinkedIn	The benefit and the risk are the same number
Eightfold AI	A 0 to 5 match score per applicant	2026 class action: hidden scores as consumer reports (FCRA)	Scoring people creates legal duties
HireVue	Video interviews at scale	Facial analysis dropped (2021); 2025 ACLU complaint, contested	Accessibility is a bias problem too

Allegations and complaints, not findings. Chosen for what each teaches, not market share.

How the Filter Gets It Wrong

Formatting

A parser misreads columns, tables, or graphics and your experience goes missing

Evidence is mostly vendor marketing; keep it simple anyway

Your name

Name alone has shifted rankings in audits of LLM screeners

UW study: white-associated names preferred about 85% of the time

History

A model trained on past hires learns past bias

Amazon dropped its recruiting model in 2018

BENEFIT

Consistent screening is possible in principle.

RISK

A bad algorithm can filter you on format, name, or inherited bias.

LLMs for Resume Writing

WHAT IT IS GOOD AT

- Catching errors and inconsistencies
- Formatting for parsers and screeners
- Tightening wording to your own facts
- Tailoring keywords to a posting

WHAT GOES WRONG

- Everyone's resume looks the same
- Flowery padding replaces facts
- Embellishment, intended or not: "CVS source control" from a new grad
- PII in a hosted prompt

BENEFIT

Cleaner, consistent, parser-friendly.

RISK

It reads like everyone else's, and can invent skills you do not have.

LLMs for Open Source Intel

WHAT YOU CAN LEARN

- What the company actually does and sells
- What the job posting translates to in daily work
- Where the team and product are heading
- Questions worth asking in the interview

THE MIRROR

- Recruiters and interviewers can run the same queries on you
- Your repos, talks, and posts are all inputs
- Models state wrong things confidently

BENEFIT

You learn the company before they learn you.

RISK

They can look you up the same way, and the model may be wrong.

OSINT on the Company

WARN Act notices

advance warning of mass layoffs

SEC EDGAR

10-K risk factors for public companies

layoffs.fyi

tech layoff tracker

News

recent coverage and stability

LinkedIn headcount trend

growing or shrinking?

Verify every claim the LLM gives you against the source.

LLMs for Mock Interviews

HOW

- Paste the job posting; ask for the interview
- Ask for follow-ups until you stall
- Better: a friend drives the LLM and asks you the questions
- Answer out loud, then review

RISK:

NONE.

JUST DO THIS!

(with a good model)

BENEFIT

Find your weak areas and practice your answers.

RISK

None, apart from time spent.

Mock Interviews

Interviews are a breadth and depth test.

BREADTH

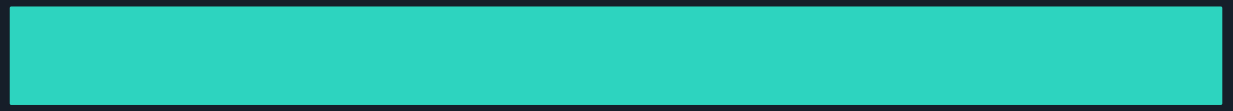

Wide and shallow: many topics from the posting

DEPTH

Narrow and deep: five follow-ups on one line

Never put anything on your resume you can't speak to at length.

Tactics

For everyone

- Solve problems
- Show evidence
- Network
- Be honest
- Aim slightly underqualified
- Tun the interview as intel
- Learn from rejection
- Keep going

Need a job, a bridge job, better pay, or a first job

- Cater to the filters
- Use an LLM to draft and revise
- Consider recruitment companies
  - Insentivised to find you a job you'll actually fit
  - You aren't obligated to say yes to any of them
  - Often get good feedback
  - MAY get a desperate employer position

More job satisfaction or a dream job

- Consider the malicious resume tactics
- Get advice from existing employees on skills needed
- Punch the time-card and keep the resume open until you land what you want

Your Interview Goal: Sell You Can Solve Their Problems

Their problem → **How you solve it** → **The risk that comes with it, and how you manage it**

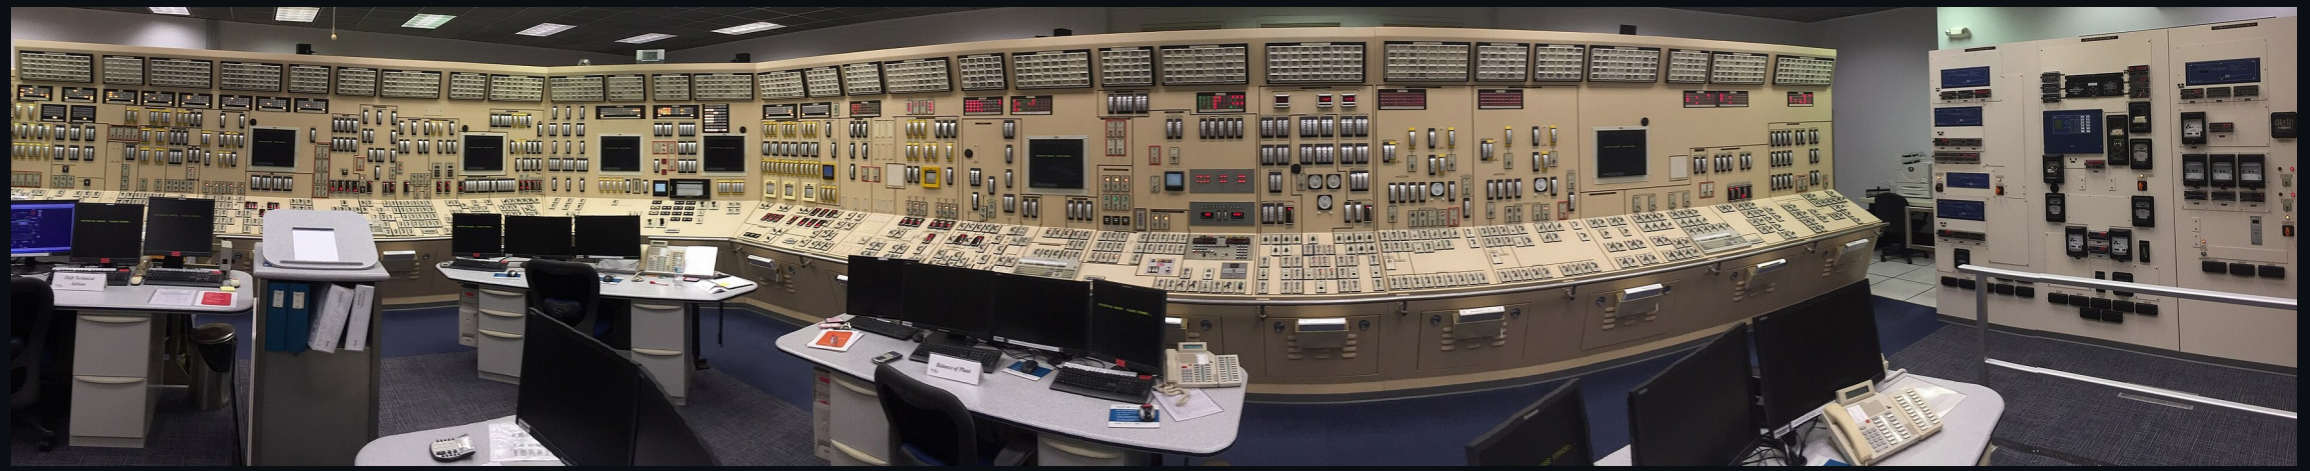

Photo: RadioFan, CC BY-SA 4.0

Vibe-coded Neopets clone

Fail forward. Worst case: a bad launch, maybe a firing.

ICS valve software for a nuclear plant

No failing forward. Executives testify before Congress.

Keep Your Resume Open

Keeping your resume open is free.

Unless the dream job is already in hand, the option costs nothing.

Market price

Learn what your skills are worth

Options

An offer you decline is still leverage

Practice

Interviews are a skill that decays

Show Tangible Evidence

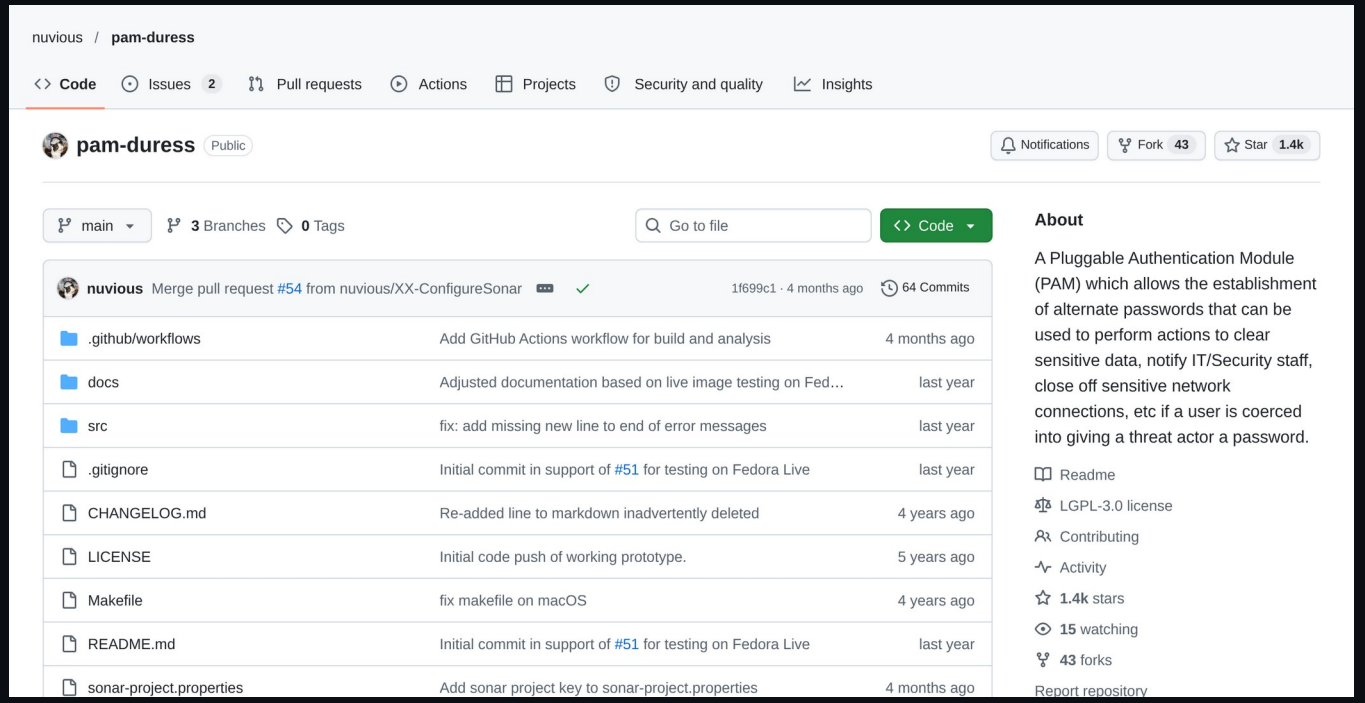

▲ PAM Duress – Alternate passwords for panic situations (github.com/nuvious)

763 points by xanthine on Aug 22, 2021 | hide | past | favorite | 343 comments

HACKER NEWS POINTS, 2021

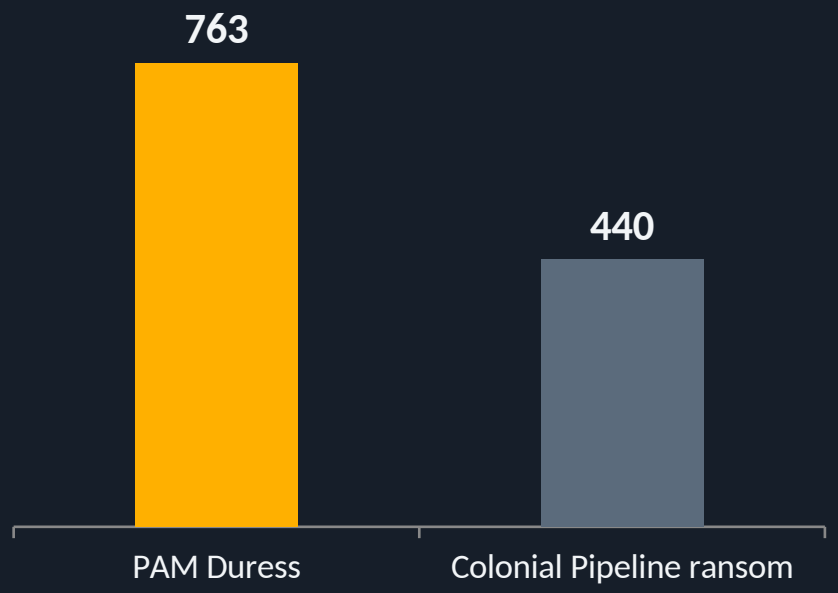

763 points, 343 comments, and an LWN write-up two days later.

Screenshots: github.com/nuvious/pam-duress and news.ycombinator.com/item?id=28267975

NETWORKING

Go where the people are.

None of it goes through an LLM screener.

- Hiring events
- Conventions
- Conferences

You are at one right now.

Conferences

You are at one right now.

LinkedIn Groups, Professional Orgs, Cold-Messaging

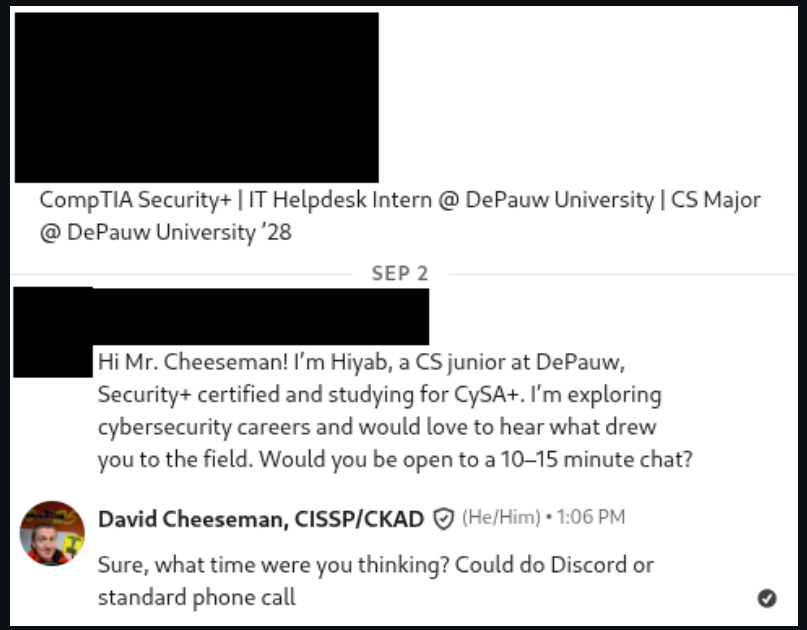

Non-Technical Groups: MAGFEST

Photo: peyj_turner, CC BY-SA 2.0 (MAGFest 2016)

How I got Hired at MAGFEST.

"What kind of skills do I need to get into the gaming industy"

Non-Technical Groups: Anime Conventions

Non-Technical Groups: Etc

Meme via mc.fly on milliways.social

Always Be Honest

DON'T

- Invent skills or experience
- Guess an answer and hope
- Let the LLM add skills you do not have

DO

- Say "I don't know"
- Offer to respond with an answer later
- Then send the answer within a day

"I don't know, but I'll find out and get back to you."

Aim for Slightly Underqualified

Overqualified

Risks slowing career progression

Slightly underqualified

Challenged, learning, and still meeting expectations

Underqualified

Risks not meeting expectations

Don’t inflate your experience or skills, it just backfires on you in the long run.

The Interview Is an Intelligence Operation

Both sides are collecting intel.

Make-or-break questions

Decide them before you walk in. Be willing to end the process yourself.

The question that does three jobs

"If I get the opportunity to work with you, what problems would I be expected to solve, and how does the team use LLMs day to day?"

Project maturity signals

New, mothballed, short staffed, behind schedule. "Did this role exist before?" hints at turnover.

Rejection Is Bidirectional Information

Ask for feedback every time (contract companies often provide this). Ending the process yourself is a rejection in the other direction, and the system working.

Don't Give Up

Keep trying.

The people who give up leave the rest of the field smaller.

Each rejection

tells you where you are not competitive yet

Each skill gap

is a thing you can go and learn

Each round

makes the next interview easier

Malicious Resume Crafting

	Hidden prompt injection	Breaking norms for humans
What it is	Invisible text telling the screener you're the best candidate	A resume that's distinctive in a way a human notices AND DELIBERATELY GETS MACHINE FILTERED!
Who it targets	The LLM	The person reading it
Payoff	Might slip past a naive screener	You're memorable to the people who matter (engineers, security professionals, etc.)
Detection	Researchers and vendors already build detectors	Nothing to detect
If caught	You're flagged as adversarial	You're remembered
Risk response	Gamble	Accept: you deliberately lose some LLM-heavy shops

Conclusion

WHAT CHANGED

- How we work: everyone is a risk manager and a mini project lead
- How employers find candidates: LLMs sift the pile before a human reads it
- How you prepare: LLMs edit your resume, research the company, and run your mock interviews

WHAT DID NOT

Risk management: weigh what could go wrong against what you stand to gain.

Own the output. The tool will not.

Resources and Contact

Connect on LinkedIn

https://www.linkedin.com/in/dcheeseman/

Slides and every reference

[short URL]

Github

github.com/nuvious

Moderated Q&A

Questions are promoted from the queue built during the talk

42:00 TO 50:00

References

- METR, early-2025 AI and experienced open-source developers (2025-07-10) and 2026-02-24 update: metr.org
- DORA 2024 (Google Cloud blog, 2024-10-22) and DORA 2025 report: dora.dev/research/2025/dora-report
- GitClear, AI Assistant Code Quality 2025: gitclear.com/ai_assistant_code_quality_2025_research
- Perry et al., Do Users Write More Insecure Code with AI Assistants?: arxiv.org/abs/2211.03622; Veracode press release, 2025-07-30
- Evaluating AGENTS.md: arxiv.org/abs/2602.11988 (listing and summaries only)
- Anthropic multi-agent research system (2025-06-13); Cognition, Don't Build Multi-Agents (2025-06-12)
- Zheng et al.: arxiv.org/abs/2311.10054; MAST: arxiv.org/abs/2503.13657
- Anthropic, AI assistance and coding skills (2026-01-29); Lee et al., CHI 2025 (Microsoft Research)
- Greenhouse 2025 reports and Fortune interview with CEO Daniel Chait; UW News on LLM resume name bias
- DFARS 252.204-7012; NIST SP 800-171 Rev 2; DOJ release on the Raytheon/RTX/Nightwing settlement (2025-05-01)

Image Credits

Image	Author	License
Shearon Harris Nuclear Power Plant control room.jpg	RadioFan	CC BY-SA 4.0
Foot Lettuce (cropped).jpg	Anonymous 4chan post, 2012, via Know Your Meme	Unlicensed meme image used as commentary
Camera filter kitten (cropped frame).jpg	394th District Court of Texas livestream, via Global News	Unlicensed frame used as commentary
USB insertion meme (Drake template).jpg	Unknown, via Know Your Meme (USB Insertion Issues)	Unlicensed meme image used as commentary
MAGFest 2016-02-19 12.11.19 HDR-2 (24947952440).jpg	peyj_turner	CC BY-SA 2.0
Run vs Run as administrator meme.png	Unknown, via mc.fly on milliways.social	Unlicensed meme image used as commentary

The screenshots on the Show Tangible Evidence slide are of public pages (GitHub, Hacker News). All other visuals are native shapes, charts, and icons.